

Course Introduction: Operating Systems

Mainack Mondal and Saptarshi Ghosh

CS30002

Spring 2021-22



Today's class

- Resources and notes
- Introduction to operating systems

Instructors



- **Mainack Mondal**
 - Research interests: system security, large scale system measurement, privacy
 - Office: CSE 316
 - Will also be in your lab

Instructors



- Saptarshi Ghosh
 - Research interests: Social Networks, ML, NLP, IR
 - Office: CSE 207
 - Will also be in your lab

Teaching assistants

- Soham Poddar (PhD)
- Abhishek Mukherjee (M.Tech.)
- Kirti Agarwal (M.Tech.)
- Adhikansh Singh (Dual)
- Ankit Bagde (Dual)
- M Koushik Raj (Dual)
- Praagy Rastogi (Dual)
- Shrey Shrivastava (Dual)
- Shivam Jha (Dual)
- Udit Dharmin Desai (Dual)
- Sanket Meshram (Dual)

Requirements

- Using computers and application software

Requirements

- Using computers and application software
- How to write, compile and run C-programs

Requirements

- Using computers and application software
- How to write, compile and run C-programs
- Computer Organization and architecture
(prerequisite)

Website / Books

- Website:
 - <https://cse.iitkgp.ac.in/~mainack/courses/2021-spring/OS-course/index.html>
- Textbooks / References:
 - Operating Systems Concepts, 9th ed. - A. Silverschatz, P.V. Galvin, and G. Gagne. Wiley
 - Some recent topics from research papers

Course logistics

- There will be 2 or 3 exams and some assignments
 - Details and mode will be announced later

Why should you attend class?

- Aside from the fact that slides are not complete and exam questions will come from class coverage?

Why should you attend class?

- Aside from the fact that slides are not complete and exam questions will come from class coverage?
- Doubt clearing sessions

Last but not the least

- Ask questions in the class
 - You need to know how Operating Systems (OS) as computer scientist
 - It is best done via class - interaction

Last but not the least

- Ask questions in the class
 - You need to know how Operating Systems (OS) as computer scientist
 - It is best done via class - interaction
- Do lab assignments & practice problems religiously
 - This is a systems course
 - There is no other way of learning how systems work than actually trying to build systems

Today's class

- Resources and notes
- Introduction to operating systems

Today's class

- Resources and notes
- Introduction to operating systems (OS)
 - What is an OS
 - What are the goals of an OS
 - Under the hood: the structure of OS
 - How does OS work?

What is an Operating System?

Let's start with enumerating the actors

Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

gcc, gedit, notepad, firefox, chrome, excel, ...

Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

gcc, gedit, notepad, firefox, chrome, excel, ...

These applications ultimately need to use hardware

processor, ram, display, mouse, keyboard

Let's start with enumerating the actors

Users

might be one or multiple (using servers)

Logs in and start applications

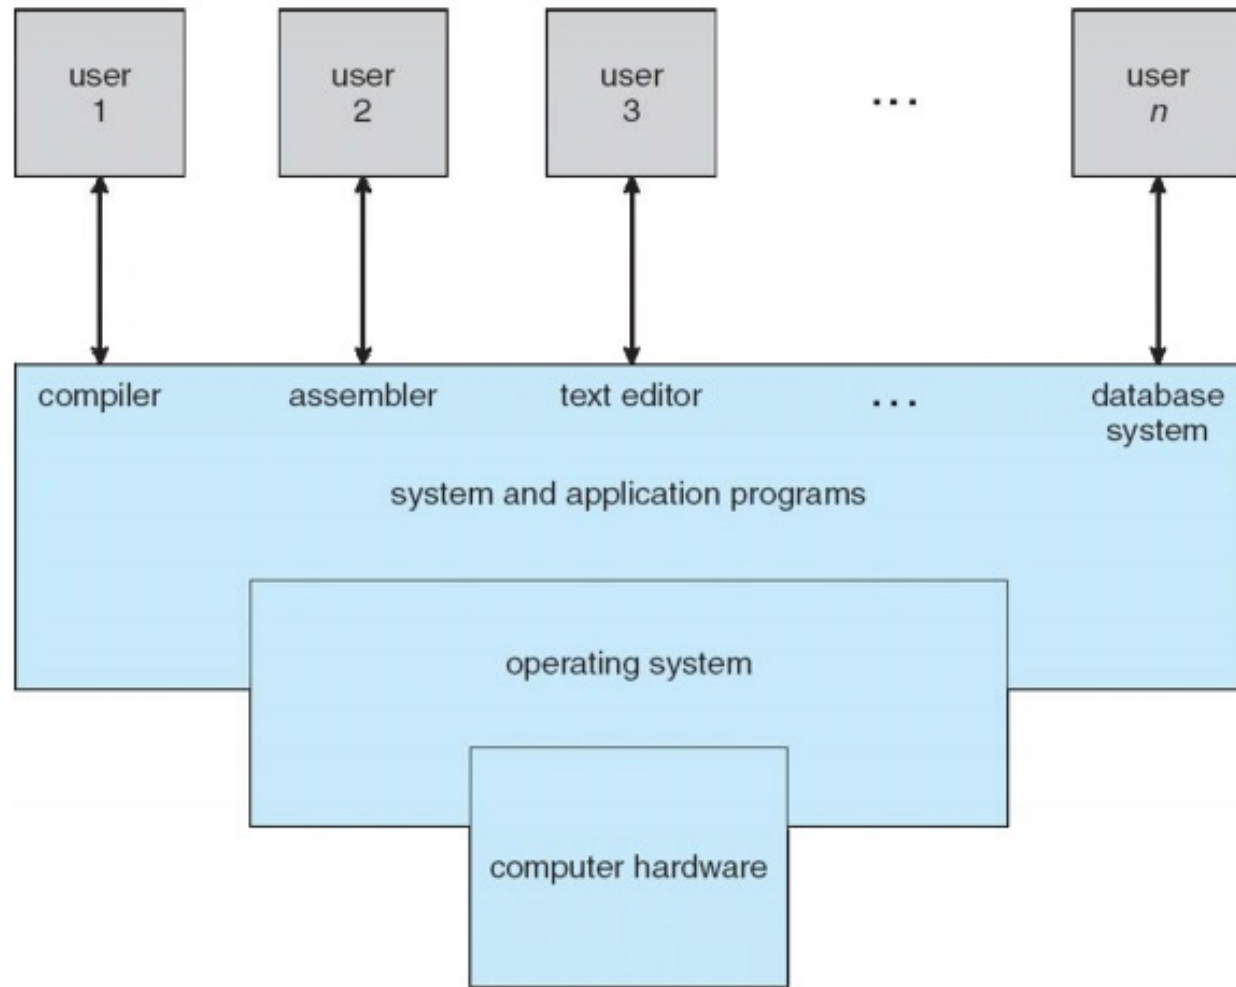
gcc, gedit, notepad, firefox, chrome, excel, ...

These applications ultimately need to use hardware

processor, ram, display, mouse, keyboard

**Operating system seats right in between applications
and hardware**

Putting it all together



OS: Definition

A program that acts as an intermediary between users of a computer system and hardware

Today's class

- Resources and notes
- Introduction to operating systems (OS)
 - What is an OS
 - What are the goals of an OS
 - Under the hood: the structure of OS
 - How does OS work?

Goals of an Operating System

Primary goals

- Make the computer system convenient and easy to use (for the users)
- Ensure efficient utilization of resources (processor time, printer, RAM, ...)
 - Controls and co-ordinates the use of hardware resources among multiple users and applications

These goals might be conflicting
across computing systems

These goals might be conflicting across computing systems

Mini computers
and mainframes
(1970 - 80)

Best utilization
of resources

Cost: user
experience
(often no
immediate
result)

These goals might be conflicting across computing systems

Mini computers
and mainframes
(1970 - 80)

Workstations /
terminals
(1980 - now)

Best utilization
of resources

Fast response
time

Cost: user
experience
(often no
immediate
result)

Cost: Resource
utilization is
suboptimal

These goals might be conflicting across computing systems

Mini computers
and mainframes
(1970 - 80)

Workstations /
terminals
(1980 - now)

Laptops

Best utilization
of resources

Fast response
time

Flexible UI, easy
to use

Cost: user
experience
(often no
immediate
result)

Cost: Resource
utilization is
suboptimal

Cost: single user
system

These goals might be conflicting across computing systems

Mini computers
and mainframes
(1970 - 80)

Workstations /
terminals
(1980 - now)

Laptops

Mobile systems

Best utilization
of resources

Fast response
time

Flexible UI, easy
to use

Optimized usability
& battery life

Cost: user
experience
(often no
immediate
result)

Cost: Resource
utilization is
suboptimal

Cost: single user
system

Cost: Severely
constrained
resources (your
android cannot
handle running full
fledged game and
browsing)

These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!

These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!

Application B: I need 2 GB of memory, now!

These goals might be conflicting across applications

Application A: I need 3 GB of memory, now!

Application B: I need 2 GB of memory, now!

Hardware: oops, total RAM is 4 GB

OS: goal-oriented definition

- A resource allocator
 - Manages all resources (processor time, RAM, display, ...) to ensure they are shared in an **efficient** and **fair** manner
 - Decides which application gets how much resources and when
- A control program
 - Controls execution of other programs / applications to prevent errors and improve usability (e.g., by giving faster response to users)
 - A faulty application should not disrupt other applications (or the OS itself)

OS: goal-oriented definition (contd.)

- Enabler of communication / coordination
 - One application may need to communicate to others, or share data / state / ...
- Enabler of easier development of applications
 - Offers a set of common services for applications (e.g., for I/O) - application developer does not need to worry about specifics of devices
 - Gives an illusion of infinite resources – dedicated processor, infinite memory, ...

Today's class


- Resources and notes
- Introduction to operating systems (OS)
 - What is an OS
 - What are the goals of an OS
 - Under the hood: the structure of OS
 - How does OS work?

Typical structure of OS (unix)

- Note: By structure we simply meant control flow


Typical structure of OS (unix)

- Note: By structure we simply meant control flow




Typical structure of OS (unix)

- Note: By structure we simply meant control flow




Typical structure of OS (unix)

- Note: By structure we simply meant control flow



Typical structure of OS (unix)

- Note: By structure we simply meant control flow



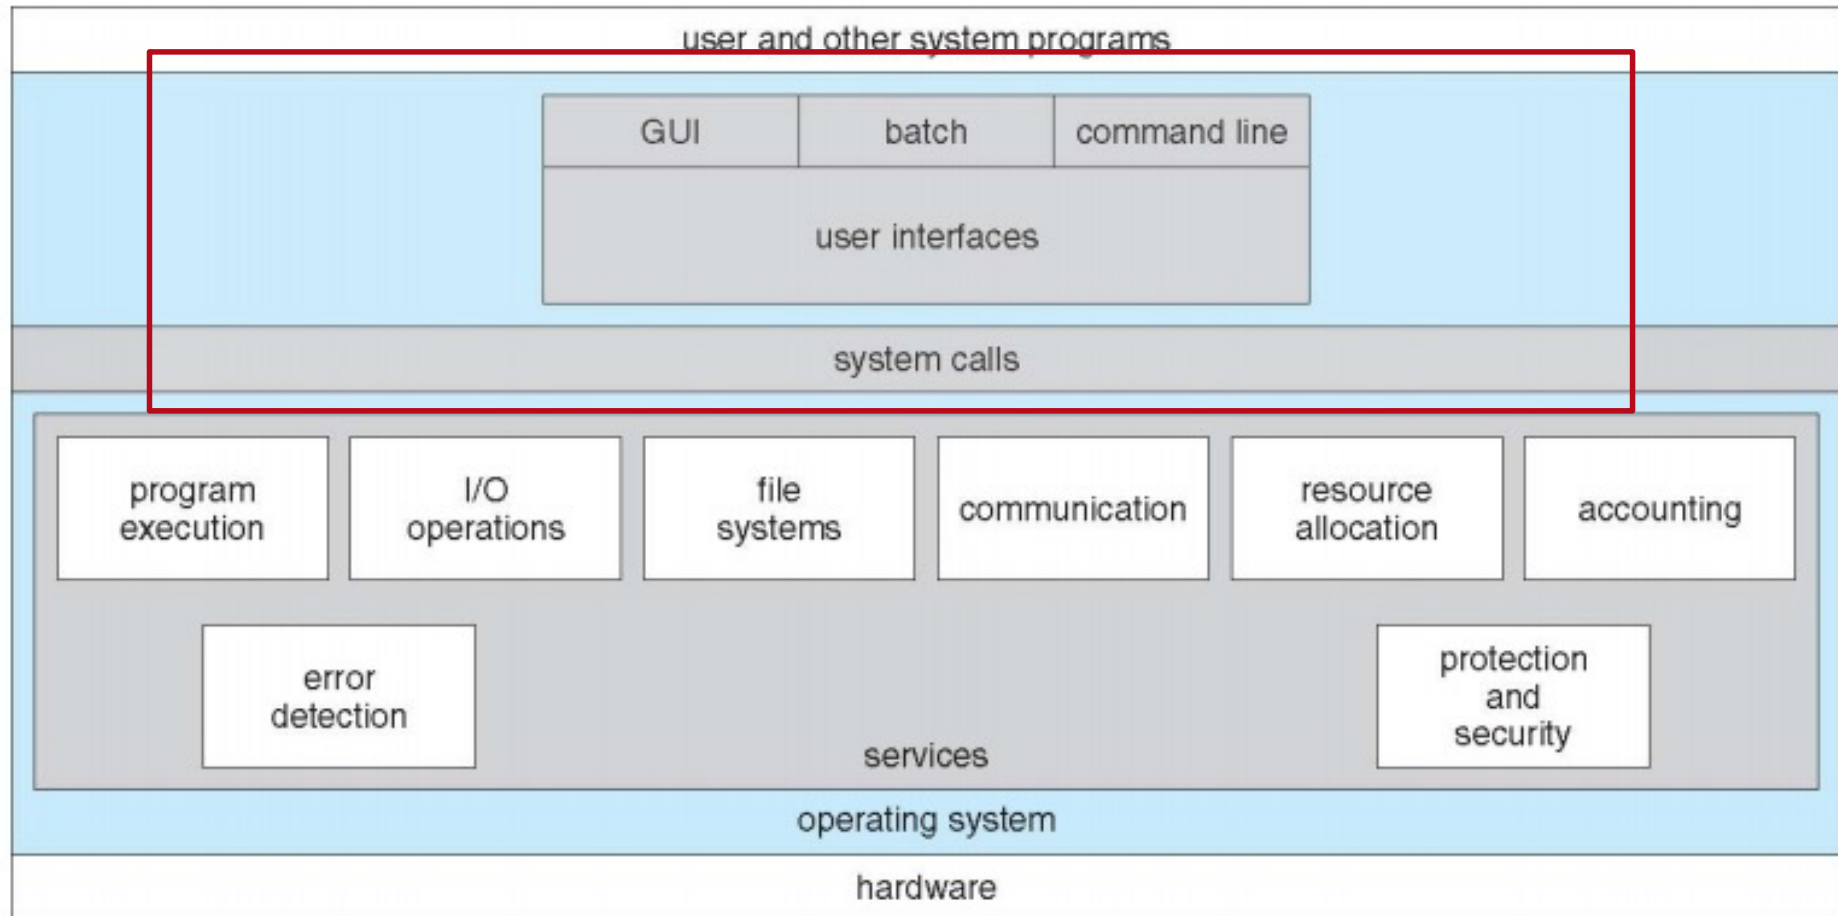
Function of Kernel

- Part of OS closest to the hardware, handles important functionalities
 - Managing memory, network, file, processes, system calls ...

Function of Kernel

- Part of OS closest to the hardware, handles important functionalities
 - Managing memory, network, file, processes, system calls ...
- Other parts of OS (e.g., the shell) and application programs can interact with kernel whenever they require these functionalities
 - E.g., need to read from keyboard (*scanf*), show something on display (*printf*), create directory (*mkdir*)

What other services?



Source: Silberschatz, Galvin and Gagne ©2013

Food for thought

- A user might ask for *printf* any time
 - Does “Kernel run all the times”?

Food for thought

- A user might ask for *printf* any time
 - Does “Kernel run all the times”?
- Possible problem
 - Kernel needs resources to run

Food for thought

- A user might ask for *printf* any time
 - Does “Kernel run all the times”?
- Possible problem
 - Kernel needs resources to run
- Insight
 - Kernel “only” needs to run when any of its functionalities are required

Today's class

- Resources and notes
- Introduction to operating systems (OS)
 - What is an OS
 - What are the goals of an OS
 - Under the hood: the structure of OS
 - How does OS work?

Today's class

- Resources and notes
- Introduction to operating systems (OS)
 - What is an OS
 - What are the goals of an OS
 - Under the hood: the structure of OS
 - How does ~~OS~~ work?
Kernel


Kernel is “interrupt” driven

- Think of interrupt as the “wake up call” to kernel
 - When interrupt comes, some function in Kernel is invoked

Kernel is “interrupt” driven


- Think of interrupt as the “wake up call” to kernel
 - When interrupt comes, some function in Kernel is invoked
- More technically
 - Interrupt is a signal (instruction), generated by hardware/software
 - The interrupt in turn activates appropriate kernel routine(s) depending on specific category of the interrupt

Types of interrupt



Hardware interrupts: generated by external device (e.g., a printer)
OR internal hardware unit

Types of interrupt

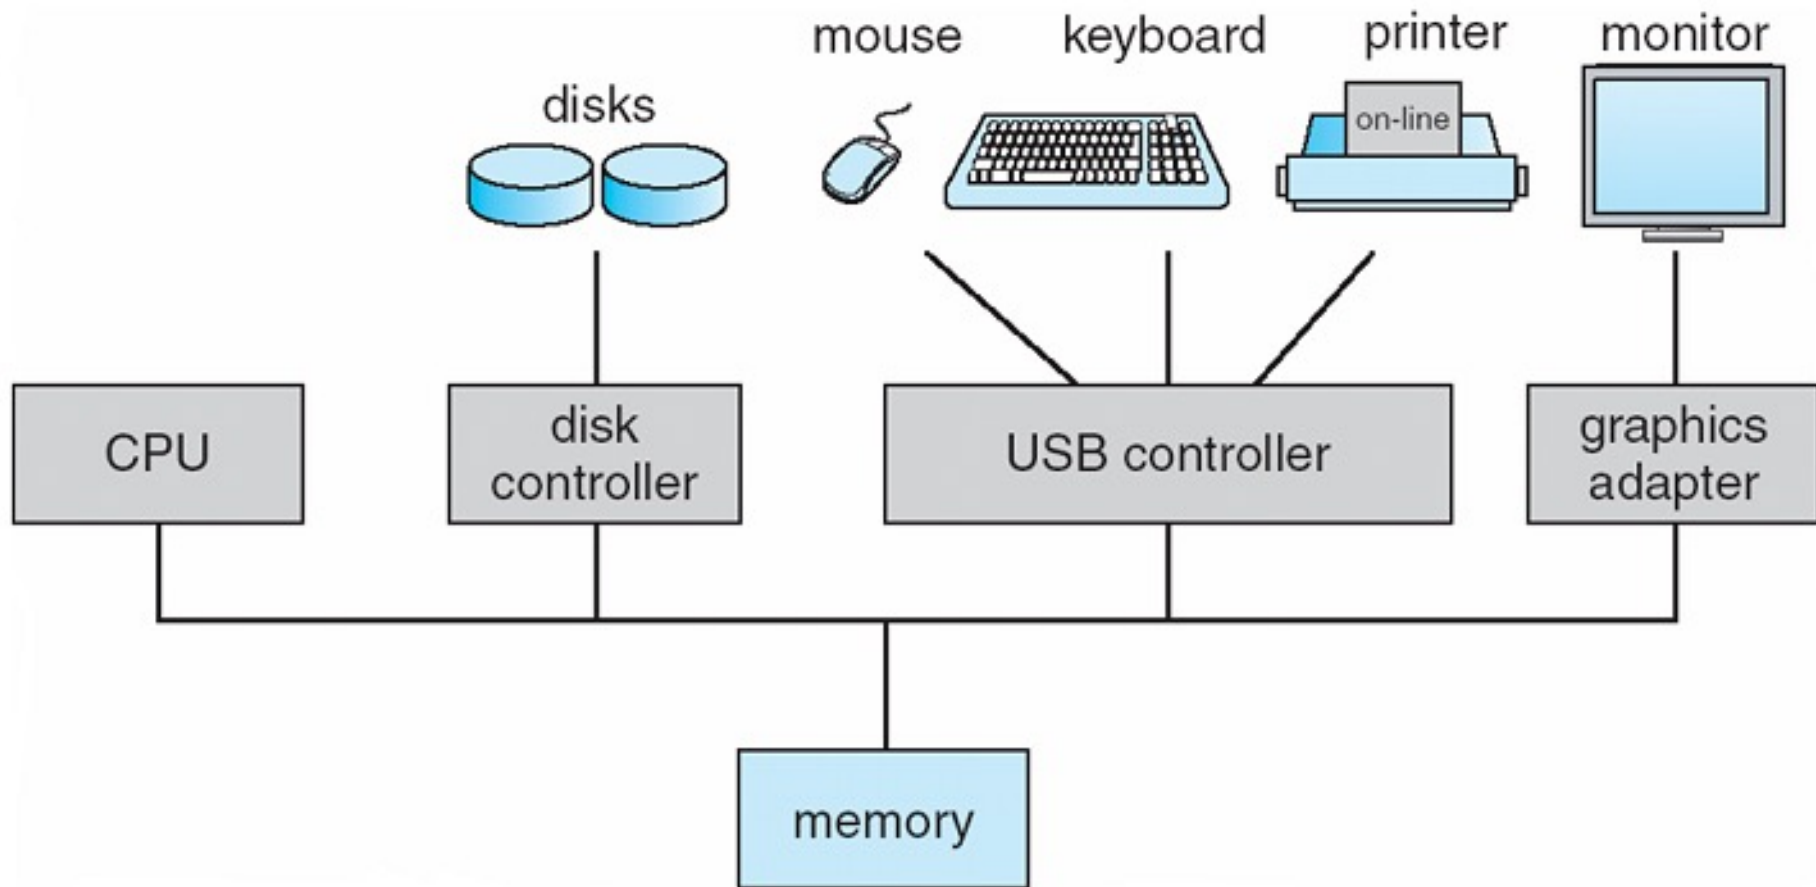


Hardware interrupts: generated by external device (e.g., a printer)
OR internal hardware unit

Software interrupts: Generated by a running program


Case study: Handling Input/Output (I/O) requests

The setup




Source: Silberschatz, Galvin and Gagne ©2013

The setup (contd.)



- I/O devices and CPU can execute in parallel
 - Each device controller is in charge of a particular device type
 - Each device controller has a local buffer
 - Data transfers between local buffer and main memory
 - Device controller sends an interrupt to the CPU to indicate I/O completion


The setup (contd.)



Question: How would an I/O request be handled in this setup?


- Each device controller has a local buffer
- Data transfers between local buffer and main memory
- Device controller sends an interrupt to the CPU to indicate I/O completion

Handling I/O in the setup




- Option 1
 - User program makes I/O request - transfers data from memory to buffer of device controller and initiates I/O
 - CPU remains idle as long as device controller handles request
 - Upon I/O completion, device controller sends interrupt, then execution of user program resumes
 - "Busy waiting" - under-utilization of CPU

Handling I/O in the setup



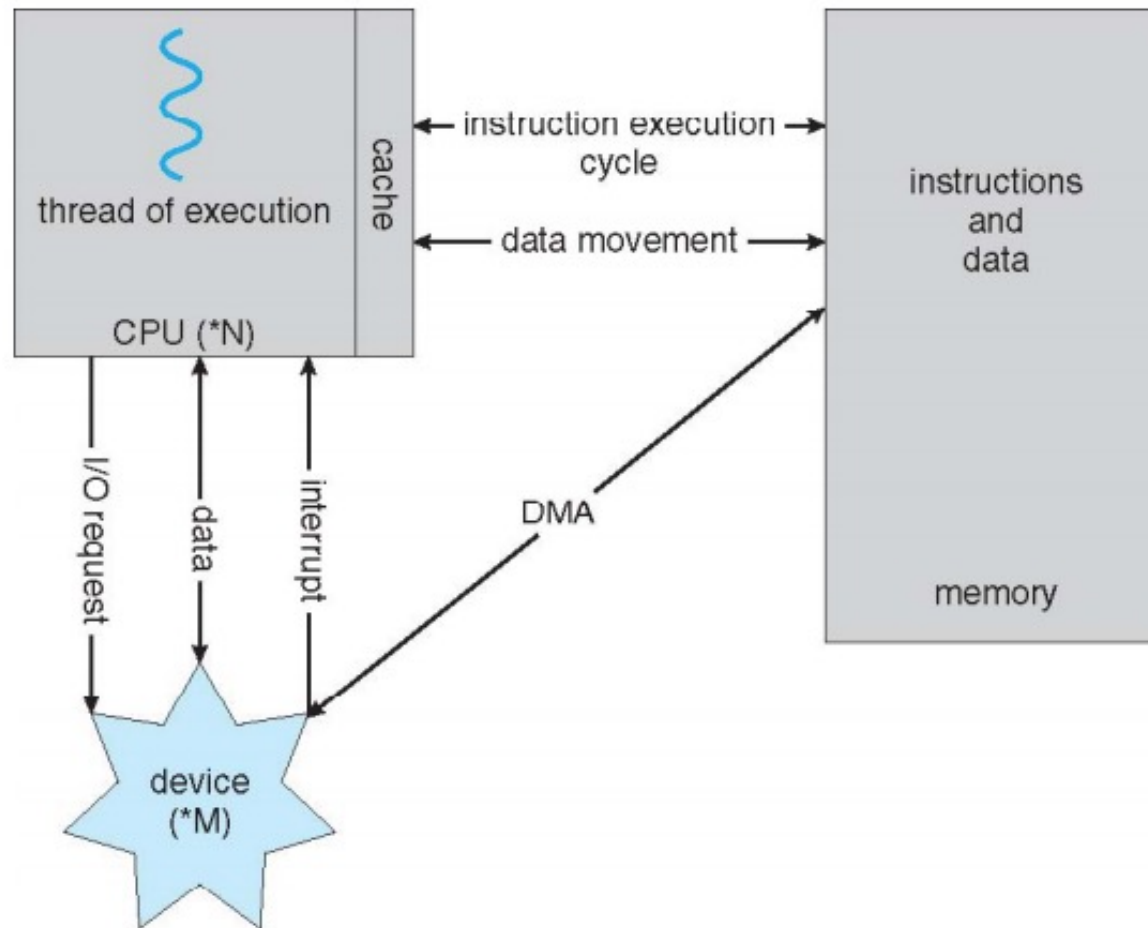
- Option 2: use OS
 - User program makes I/O request through a "system call"
 - OS transfers data from memory to buffer of device controller and initiates I/O
 - OS can allocate CPU to other programs during I/O of one program
 - Upon I/O completion, device controller sends interrupt, then OS may resume this particular user program (according to some scheduling policy)
 - "Multiprogramming" - much better utilization of CPU

Handling I/O in the setup



- Previous scheme of interrupt-driven I/O is fine for small amounts of data, but high overhead for bulk data transfer such as disk I/O
- Hence option 3 – Direct Memory Access (DMA)
 - Once OS initiates I/O, DMA manages transfer of data between memory and device controller with no intervention by the CPU

I/O handling in von Neumann architecture



Source: Silberschatz, Galvin and Gagne ©2013