

System calls

Saptarshi Ghosh and Mainack Mondal
CS39002

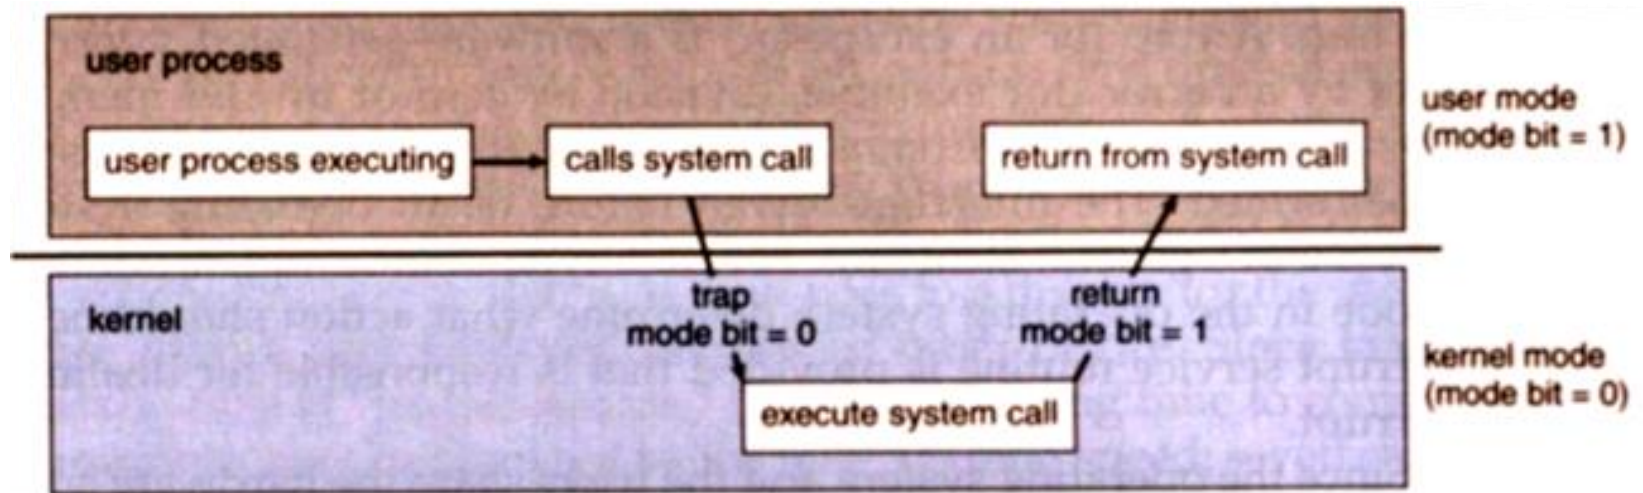
Spring 2020-21



Some slides borrowed from slides of Don Porter, Univ of North Carolina (UNC)

What are system calls?

- The mechanism used by an application program to request service from the operating system
 - So how does it work?



This lecture

- How system calls work
 - Overview of interrupts
 - Earlier days – syscalls handled similar to interrupts
 - Nowadays – how syscalls work
- Examples of common syscalls

Earlier days: interrupt

- Originally, system calls issued using “int” instruction and handled similar to interrupts
 - Let us see a brief overview of how interrupts work

Interrupts

Interrupts

- Two types of interrupts usually provided by a processor

Interrupts

- Two types of interrupts usually provided by a processor
- Synchronous – will happen every time an instruction executes with a given program state
 - Divide by zero
 - Bad pointer dereference
 - Application wants some service (syscall)

Interrupts

- Two types of interrupts usually provided by a processor
- Synchronous – will happen every time an instruction executes with a given program state
 - Divide by zero
 - Bad pointer dereference
 - Application wants some service (syscall)
- Asynchronous – caused by an external event
 - By I/O device, timer, etc

Interrupts: Intel nomenclature

- Two types of interrupts usually provided by a processor
- Synchronous – will happen every time an instruction executes with a given program state
 - Divide by zero
 - Bad pointer dereference
 - Application wants some service (syscall)
- Asynchronous – caused by an external event
 - By I/O device, timer, etc

Exceptions

Interrupts

Interrupts: Intel nomenclature

- Two types of interrupts usually provided by a processor
- Synchronous – will happen every time an instruction executes with a given program state
 - Divide by zero
 - Bad pointer dereference
 - Application wants some service (syscall)
- Asynchronous – caused by an external event
 - By I/O device, timer, etc

Exceptions

Interrupts

From programmer's perspective, handled with the same abstractions


How interrupts work

- Each interrupt / exception includes a number indicating its type (e.g., 14 is a page fault)
 - Kernel creates an array of Interrupt descriptors in memory, called **Interrupt Descriptor Table**, or IDT
 - The number is index into the IDT

How interrupts work

- Each interrupt / exception includes a number indicating its type (e.g., 14 is a page fault)
 - Kernel creates an array of Interrupt descriptors in memory, called **Interrupt Descriptor Table**, or IDT
 - The number is index into the IDT
- IDT can be anywhere in memory
- Pointed to by a special register (idtr)
- Each IDT entry j specifies things like
 - What code to execute for interrupt j (called interrupt handler)
 - Privilege level necessary to raise the interrupt (e.g., a user program cannot raise interrupt for page fault)

Example: x86 interrupt table



Software interrupts

- The "int <num>" instruction allows software to raise an interrupt
- What happens if a program issues "int <num>"?

Software interrupts

- The "int <num>" instruction allows software to raise an interrupt
- What happens if a program issues "int <num>"?
 - Control jumps to the kernel at a prescribed address (the interrupt handler)
 - Address of interrupt handler found using the IDT (index num)
 - The register state of the program is saved by the kernel
 - Interrupt handler runs and handles the interrupt
 - When handler completes, resume program

An important concept

- The state of a program's execution is succinctly and completely represented by the CPU register state
- Pause a program (so that some other program can be allocated the CPU)
 - Dump the register contents into memory
- Resume a program
 - Read the register contents back from memory to the CPU

System calls

Earlier days

- Originally, system calls issued using “int” instruction

Earlier days

- Originally, system calls issued using “int” instruction
 - The system call handler routine was just an interrupt handler

Earlier days

- Originally, system calls issued using “int” instruction
 - The system call handler routine was just an interrupt handler
 - Like interrupts, system calls used to be arranged in a table

Earlier days

- Originally, system calls issued using “int” instruction
 - The system call handler routine was just an interrupt handler
 - Like interrupts, system calls used to be arranged in a table
- Whenever a syscall (interrupt driven) came

Earlier days

- Originally, system calls issued using “int” instruction
 - The system call handler routine was just an interrupt handler
 - Like interrupts, system calls used to be arranged in a table
- Whenever a syscall (interrupt driven) came
 - Program selects which syscall it wants by placing index in a particular register (eax)

Earlier days

- Originally, system calls issued using “int” instruction
 - The system call handler routine was just an interrupt handler
 - Like interrupts, system calls used to be arranged in a table
- Whenever a syscall (interrupt driven) came
 - Program selects which syscall it wants by placing index in a particular register (eax)
 - Arguments go in other registers as per convention
 - Return value goes in eax

Earlier days: summary

- Kernel leveraged interrupts to handle system calls

However this mechanism is slow

- Today, processors are totally pipelined
 - Pipeline stalls are very expensive
 - Cache misses can cause pipeline stalls
- Recall that IDT is in memory
 - May not be in cache
 - Cache misses makes it expensive, e.g., system calls said to take almost twice as long from P3 to P4

Idea: new instruction

- What if we cache the IDT entry for a system call in a special CPU register?
 - No more cache misses for the IDT!
- What is the cost?

Idea: new instruction

- What if we cache the IDT entry for a system call in a special CPU register?
 - No more cache misses for the IDT!
- What is the cost?
 - system calls should be frequent enough to be worth the transistor budget

How to leverage the new instruction?

- There is a machine instruction (new architectures)
 - Essentially for asking your processor to perform task
 - Hardware specific
 - Can be called “syscall”, “trap”, “svc”, “swi”
 - For x86-64 architecture it is called “syscall”

Example of syscall

- syscall machine instruction takes operands
 - syscall 10 // integer number for x86
// some of these are fixed by intel
- The instruction used machine-specific registers that store syscall entry point (into kernel code) for every syscall
- Arguments, return value go in other registers as per convention

System calls in programs

Syscalls in programs

- When we are writing programs in HLL (higher level language, think C)
 - We think of syscalls in a somewhat higher level context

Characteristics of system calls

- Provide an interface to the services made available by an operating system

Characteristics of system calls

- Provide an interface to the services made available by an operating system
 - Typically executes hundreds of thousands time every second

Characteristics of system calls

- Provide an interface to the services made available by an operating system
 - Typically executes hundreds of thousands time every second
 - User application programs do not see this level of detail

Characteristics of system calls

- Provide an interface to the services made available by an operating system
 - Typically executes hundreds of thousands time every second
 - User application programs do not see this level of detail
 - Use Application programming interface (API)

Characteristics of system calls

- Provide an interface to the services made available by an operating system
 - Typically executes hundreds of thousands time every second
 - User application programs do not see this level of detail
 - Use Application programming interface (API)
 - fork, pipe, execvp etc.

Characteristics of system calls

- Provide an interface to the services made available by an operating system
 - Typically executes hundreds of thousands time every second
 - User application programs do not see this level of detail
 - Use Application programming interface (API)
 - fork, pipe, execvp etc.
 - These APIs are also loosely termed as system calls

Example of a system call API

- There are no `fopen`, `fgets`, `printf`, and `fclose` system calls in the Linux kernel but `open`, `read`, `write`, and `close`
- <http://man7.org/linux/man-pages/man2/read.2.html>

NAME

[top](#)

`read` - read from a file descriptor

SYNOPSIS

[top](#)

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

Example of a system call API

- <http://man7.org/linux/man-pages/man2/read.2.html>

NAME [top](#)

read - read from a file descriptor

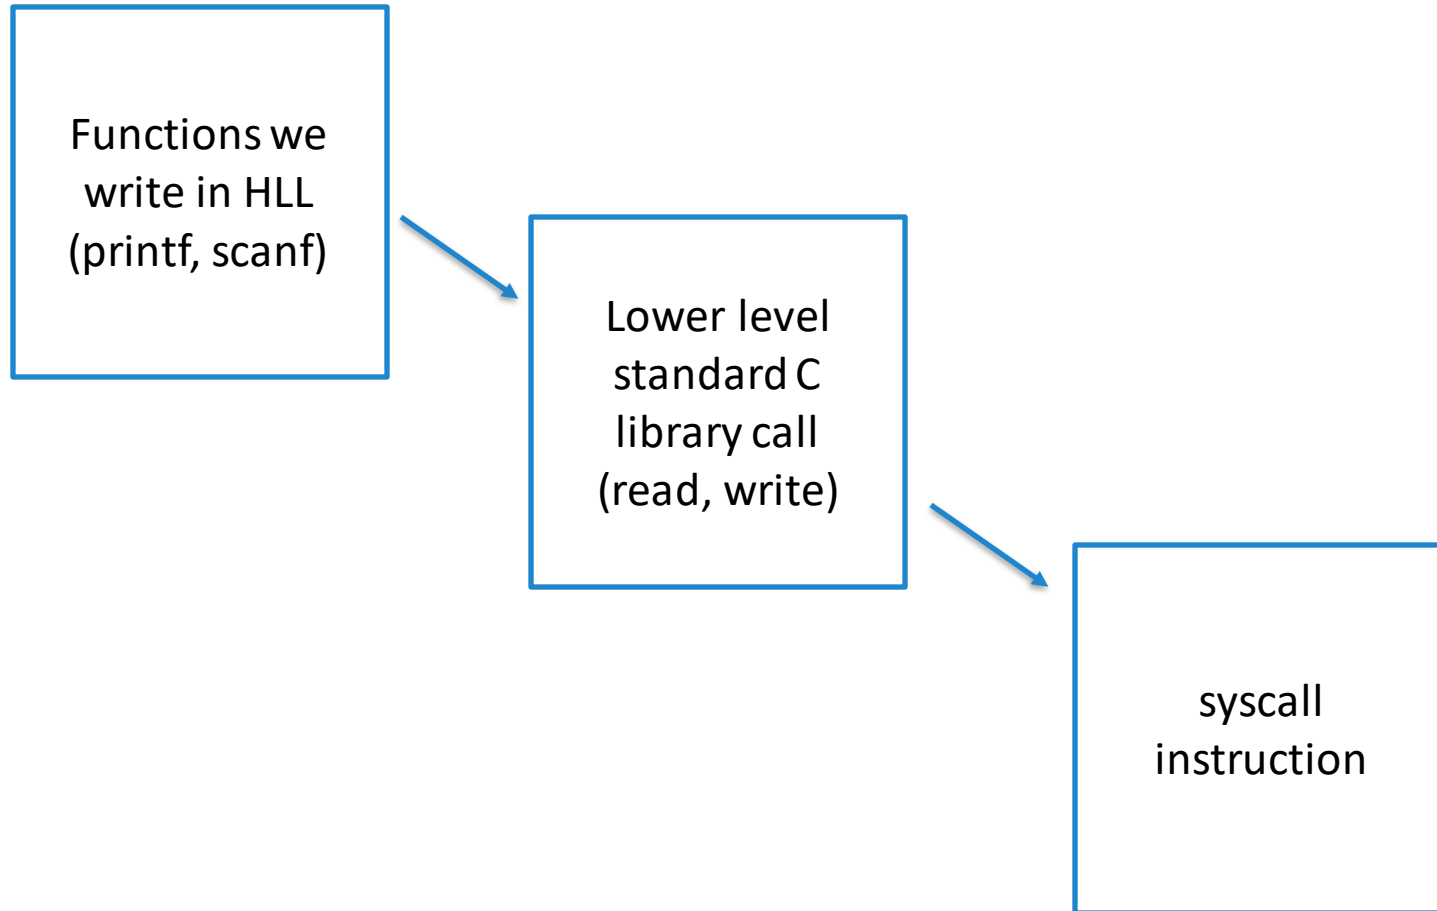
SYNOPSIS [top](#)

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

- What are these function parameters?
- What is the return values?

The workflow of a syscall



Types of system calls

(From silberschatz's slides)

Types of System Calls

- **Process control (e.g., fork(), exit(), wait())**
 - create process, terminate process (fork, exit)
 - end, abort
 - load, execute
 - get process attributes, set process attributes
 - wait for time
 - wait event, signal event
 - allocate and free memory
 - Dump memory if error
 - **Debugger** for determining **bugs, single step** execution
 - **Locks** for managing access to share data between processes

Types of System Calls (Cont.)

- **File management (e.g., open(), close(), read(), write())**
 - create file, delete file
 - open, close file
 - read, write, reposition
 - get and set file attributes
- **Device management (e.g., ioctl(), read(), write())**
 - request device, release device
 - read, write, reposition
 - get device attributes, set device attributes
 - logically attach or detach devices

Types of System Calls (Cont.)

- Inter-Process Communications (e.g., pipe(), semget(), semop(), shmget(), shmcat(), shmdt(), shmctl(), signal(), kill())
 - create, delete communication connection
 - send, receive messages if **message passing model** to **host name** or **process name**
 - **Shared-memory model** create and gain access to memory regions
 - transfer status information
 - attach and detach remote devices

Types of System Calls (Cont.)

- **Protection (chmod(), chown(), umask())**
 - Control access to resources
 - Get and set permissions
 - Allow and deny user access

More about system calls

- System call numbers in linux for x86 :
 - https://github.com/torvalds/linux/blob/16f73eb02d7e1765ccab3d2018e0bd98eb93d973/arch/x86/entry/syscalls/syscall_32.tbl
- System call numbers in linux for x86 - 64:
 - https://github.com/torvalds/linux/blob/16f73eb02d7e1765ccab3d2018e0bd98eb93d973/arch/x86/entry/syscalls/syscall_64.tbl
- System call implementations in x86-64 and x86
 - <https://stackoverflow.com/questions/15168822/intel-x86-vs-x64-system-call>