

Process (contd.)

Saptarshi Ghosh and Mainack Mondal
CS39002

Spring 2020-21



So far on processes

- What is a process?
- Structure of a process in memory
- Process control block (PCB)
- Process states
- Context switch

Process queues (for scheduling)

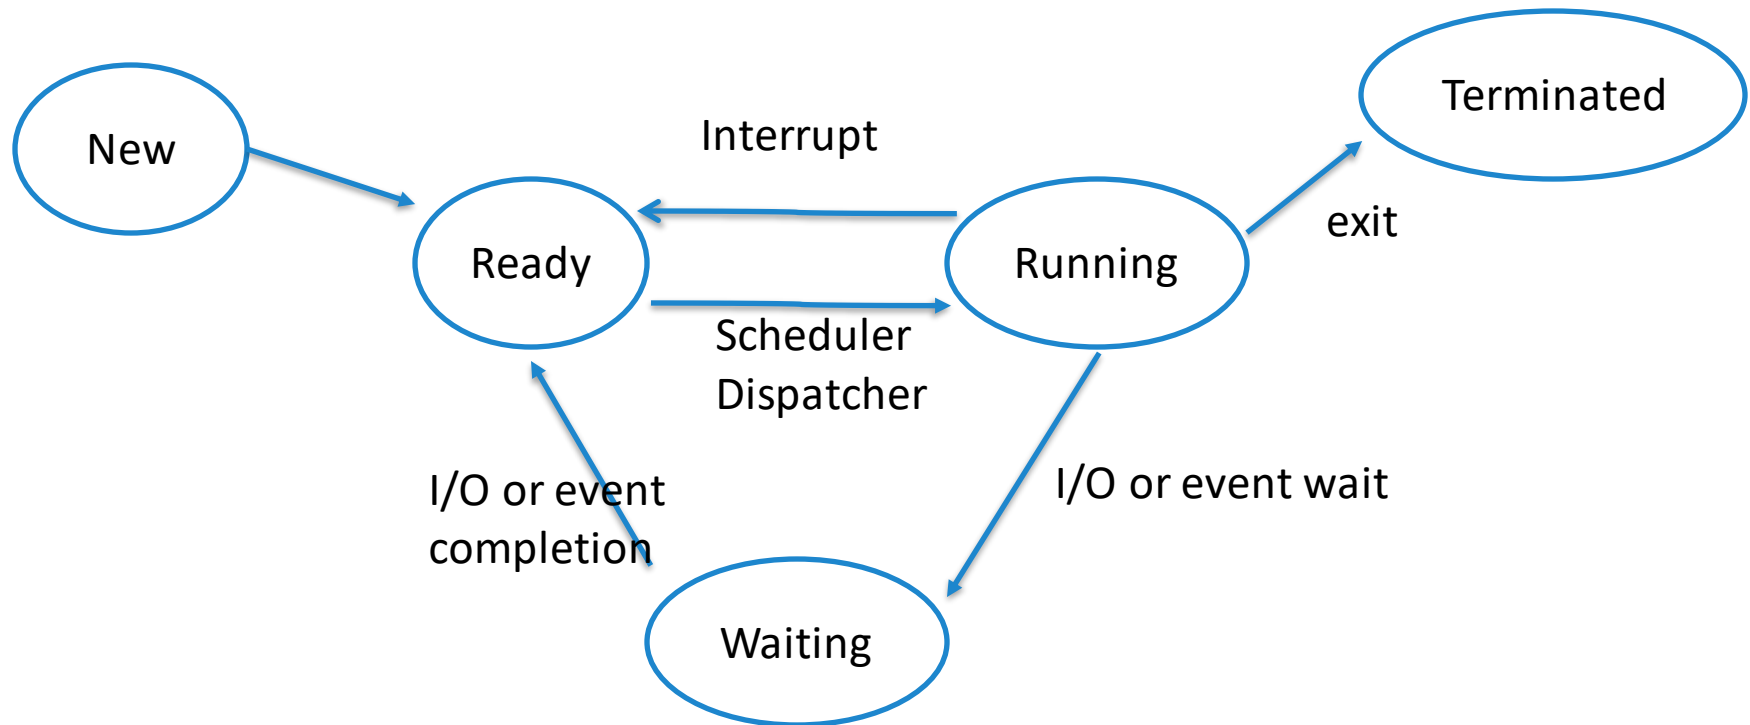
Process queues for scheduling

- Several scheduling queues exist in OS
 - A PCB is linked to one of the queues at any given time
 - The PCBs in a queue are connected as a linked list

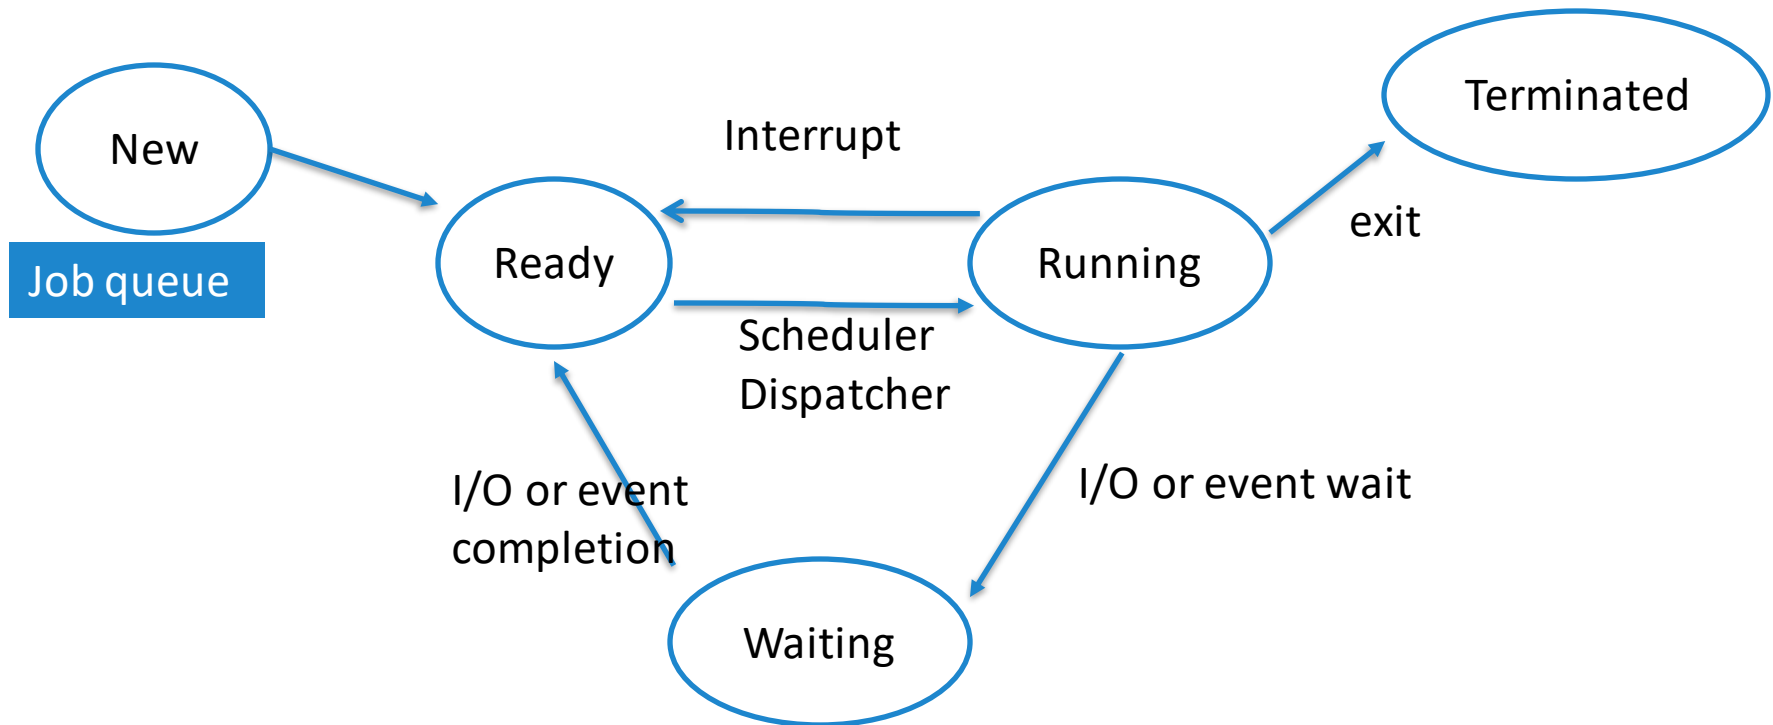
Characteristics of process queues

- Each I/O device has its own device queue
 - Contains (PCBs of) processes waiting for this device
- Each event also has its own queue
 - Contains processes waiting for this event
- Process scheduling can be represented as a queueing diagram
 - Queueing diagram represents queues, resources, flows
 - We will discuss actual scheduling algorithms later

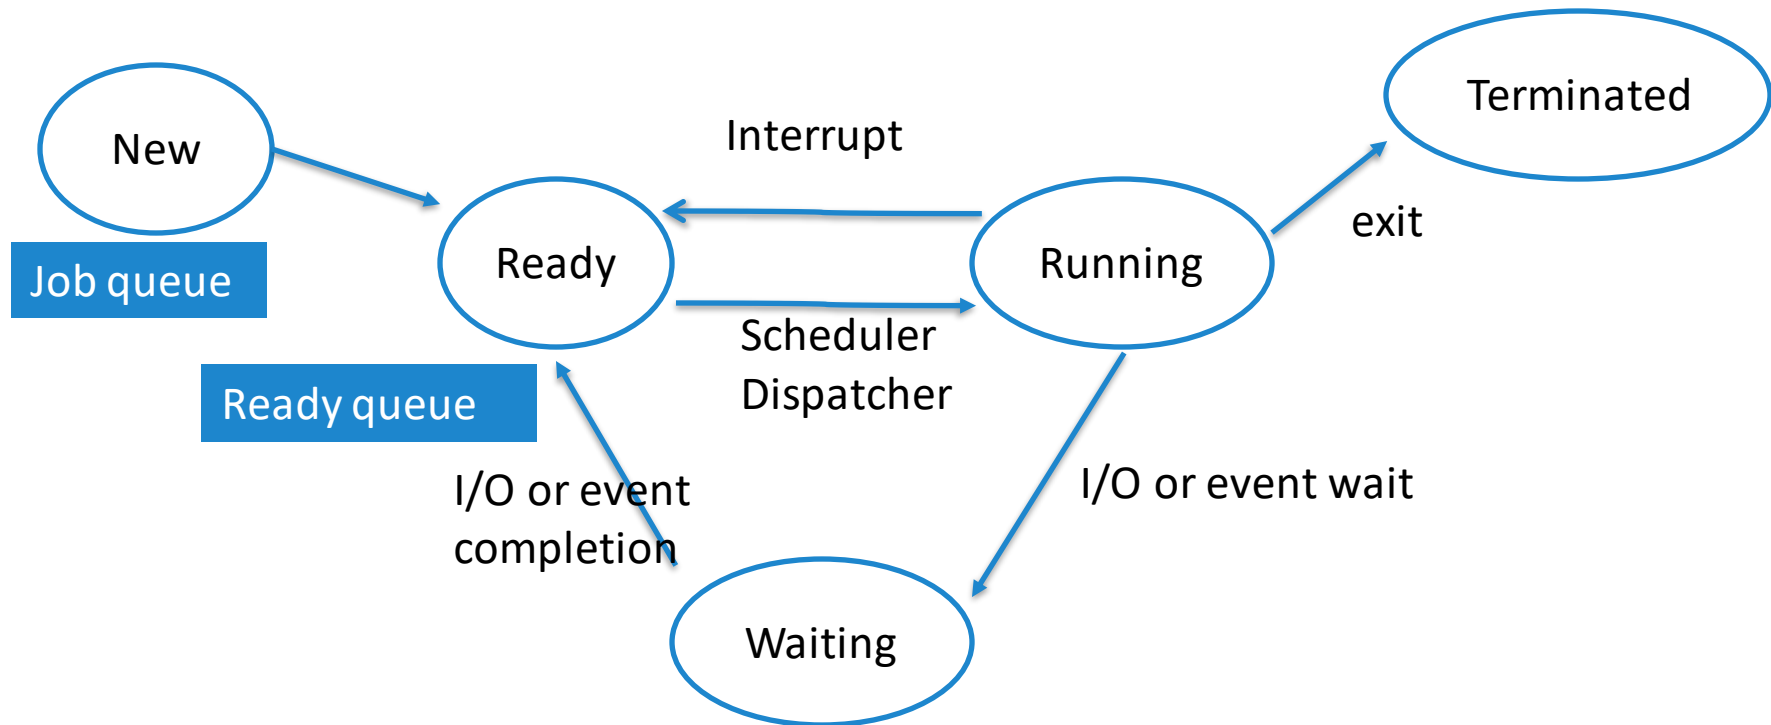
Recap: Process state diagram



Recap: Process state diagram




Recap: Process state diagram



- The **process scheduler** selects a process from the ready queue for execution on the CPU
- **Dispatcher**: The kernel process that assigns the CPU to a process selected by the scheduler

Recap: Process state diagram

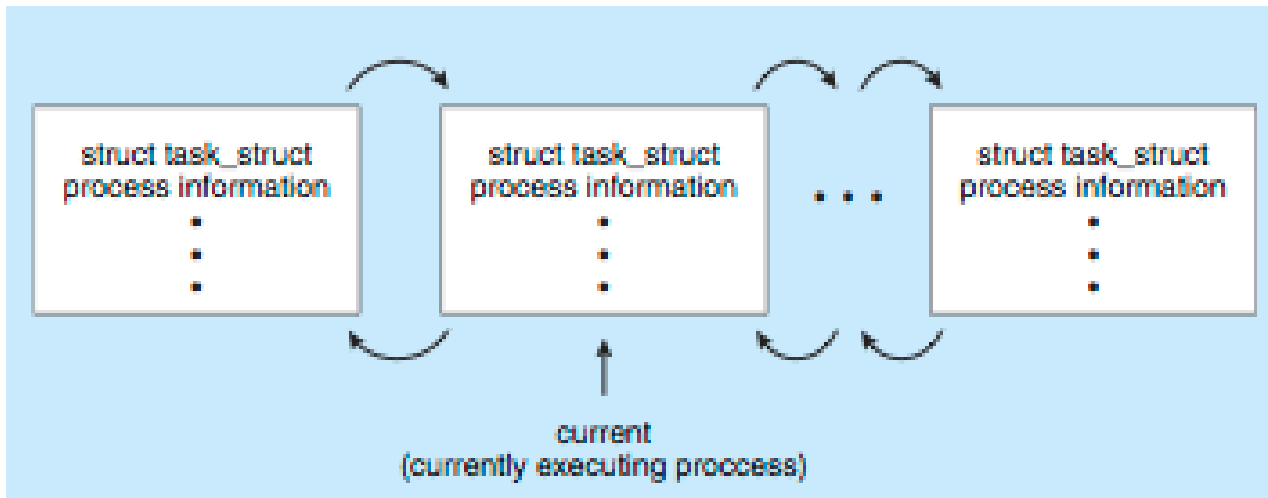


Process queue representation in Linux


Represented by the C structure `task_struct`

```
pid_t pid; // process identifier
long state; // state of the process
unsigned int time slice; // scheduling information
struct task_struct *parent; // this process's parent
struct task_struct *children; // this process's children
struct files_struct *files; // list of open files
struct mm_struct *mm; // address space of this pro
```

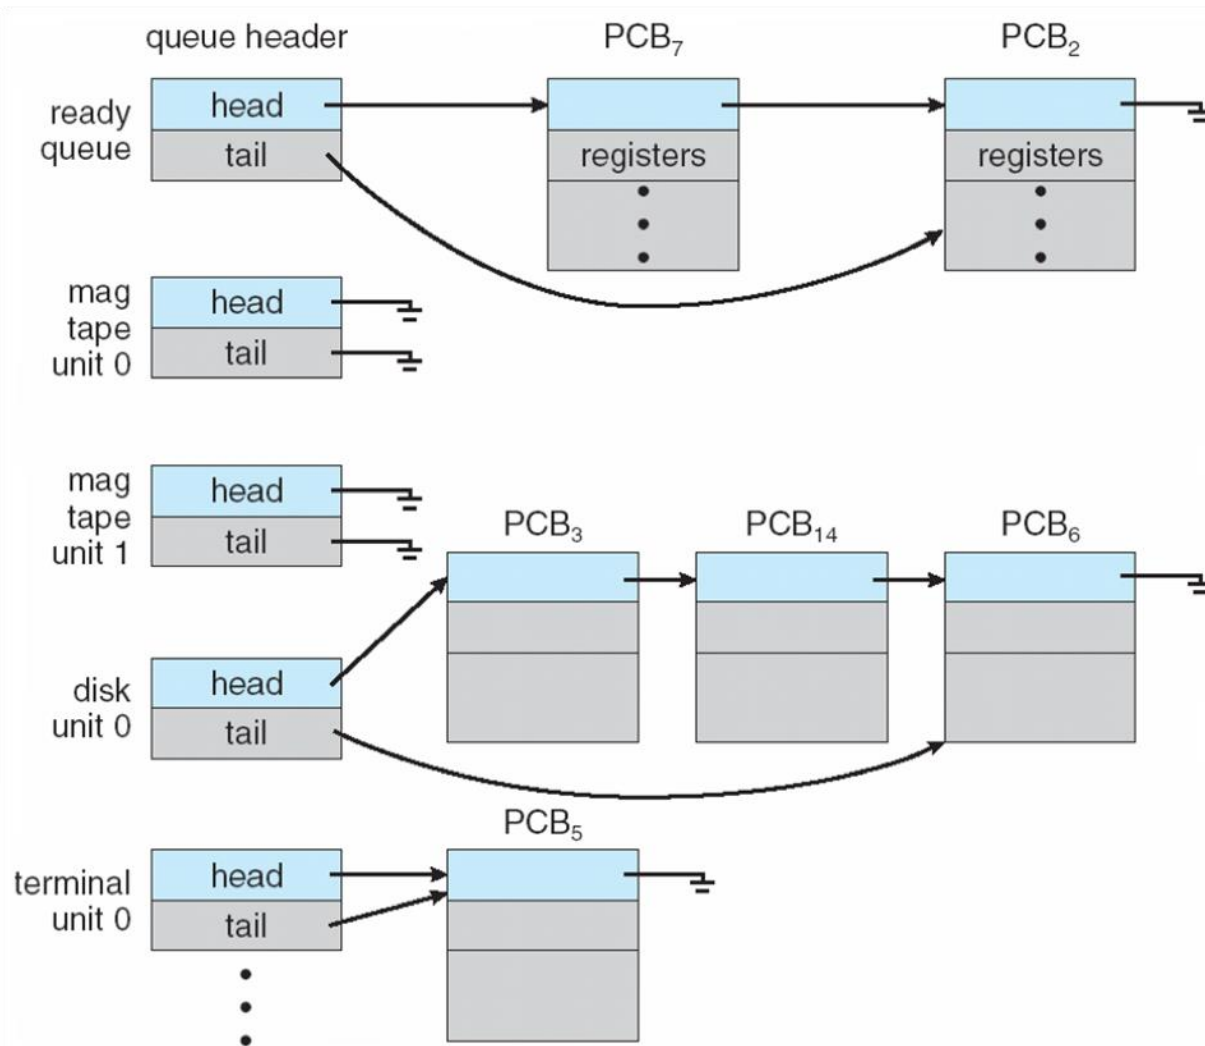
Doubly
linked list



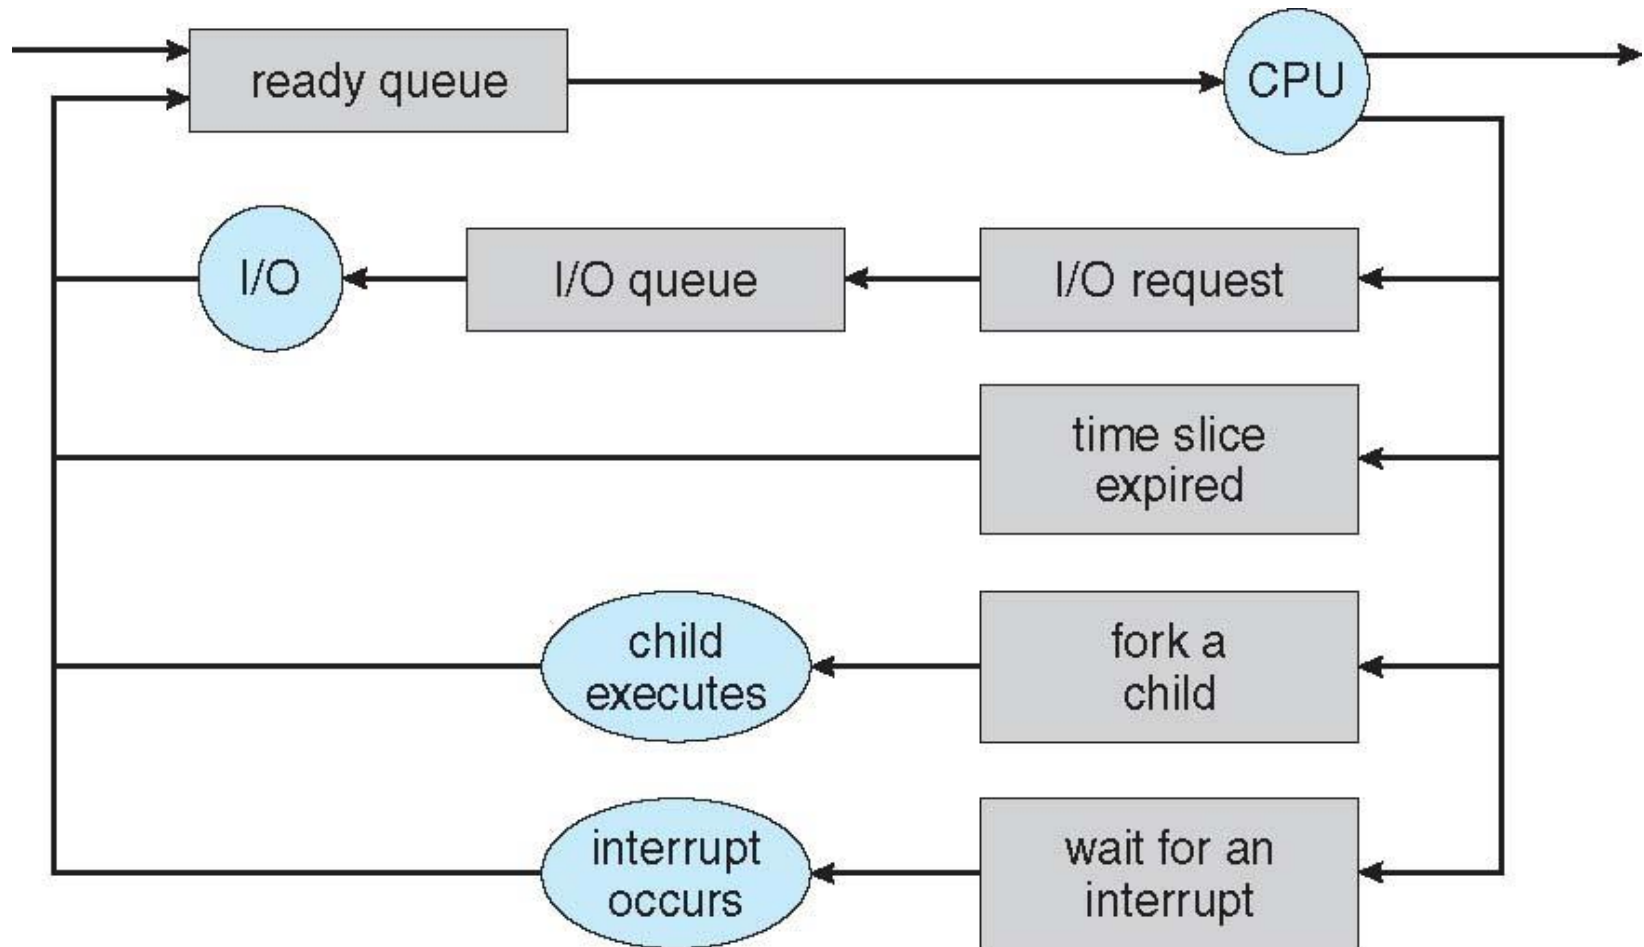
Structure of process queues



Structure of process queues



Representation of process scheduling



Operations on processes

Process creation

- During execution a process may create several new processes
 - The process that creates a new process: **parent process**
 - The new process created: **child process**
 - Parent process create children processes, which, in turn, can create other processes, forming a tree of processes
 - Each process has a unique process identifier (pid)
 - Other than the first process (init), all other processes are created by **fork() system call**

Process creation (contd.)

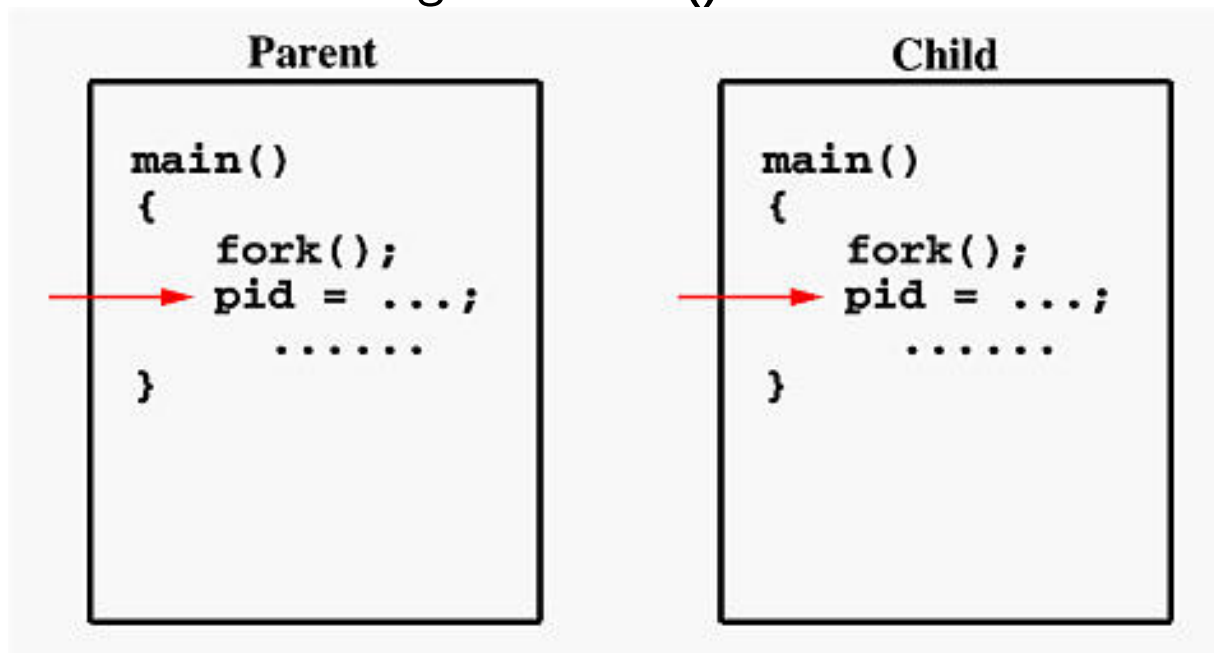
- Address space
 - Initially child is duplicate of parent
 - Child can later have a different program loaded into it
- In UNIX / Linux
 - `fork()` : creates a new process
 - `exec()`: replace new process's memory with new code

fork() system call

- If **fork()** returns a negative value, the creation of a child process was unsuccessful.
- **fork()** returns a zero to the newly created child process.
- **fork()** returns a positive value, the *process ID* of the child process, to the parent.

fork() system call

- If the call to **fork()** is executed successfully, the kernel will make an identical copy of the parent's address space, for use by the child.
- Both processes will start their execution at the next statement following the **fork()** call.




Make a child do something different than the parent – method 1

Parent

```
main()      pid = 3456
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```




Child

```
main()      pid = 0
{
    pid=fork();
    if (pid == 0)
        ChildProcess();
    else
        ParentProcess();
}

void ChildProcess()
{
    .....
}

void ParentProcess()
{
    .....
}
```



Make a child do something different than the parent – method 2

- `exec()` system call and its variants
 - Replaces the calling process (that calls `exec`) with another program
 - Calling process terminated; the specified program is loaded in the same address space and then executed
 - New process image has same process ID (pid), same file descriptors, etc
 - Several variants – `execl`, `execle`, `execvp`, `execv`, `execve`, `execvp`

Process waiting for a child

- A parent process can wait for a child to complete
 - Parent indicates to the kernel that it wants to wait, using the `wait()` system call
 - Parent execution is suspended; will be "ready" only after child terminates
 - If a child has already terminated, then the call returns immediately. Otherwise, block until a child terminates

Process waiting for a child

- A parent process can wait for a child to complete
 - Parent indicates to the kernel that it wants to wait, using the `wait()` system call
 - Parent execution is suspended; will be "ready" only after child terminates
 - If a child has already terminated, then the call returns immediately. Otherwise, block until a child terminates
- Two variants
 - `wait()` - suspends execution of the current process until one of its children terminates
 - `waitpid(pid)` - suspends execution of the current process until the child specified by *pid* terminates

Process termination

- A child process executes last statement
 - `exit()` call for deleting the process
 - Return status data to parent via `wait()` if any parent is waiting for the termination of this process
 - Deallocate the resources

Process termination

- A child process executes last statement
 - `exit()` call for deleting the process
 - Return status data to parent via `wait()` if any parent is waiting for the termination of this process
 - Deallocate the resources

Child process

.
.

`exit(2) // Exit with status code`

Parent process

`pid_t pid;`

`int status;`

.

`pid = wait (&status) // pid of
terminated child`

Process creation and exec: example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { Error condition
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { child process
        execlp("/bin/ls", "ls", NULL);
    }
    else { parent process
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Process termination: Special cases

- In some OS
 - All child must terminate when a process terminates
 - Cascading termination: All children, grandchildren etc. must be terminated when a process terminates
 - OS takes care of this cascade
- Combinations of `exit()` and `wait()`
 - If no parent is waiting, then **zombie process**
 - If parent terminated without invoking `wait()` then **orphan process**

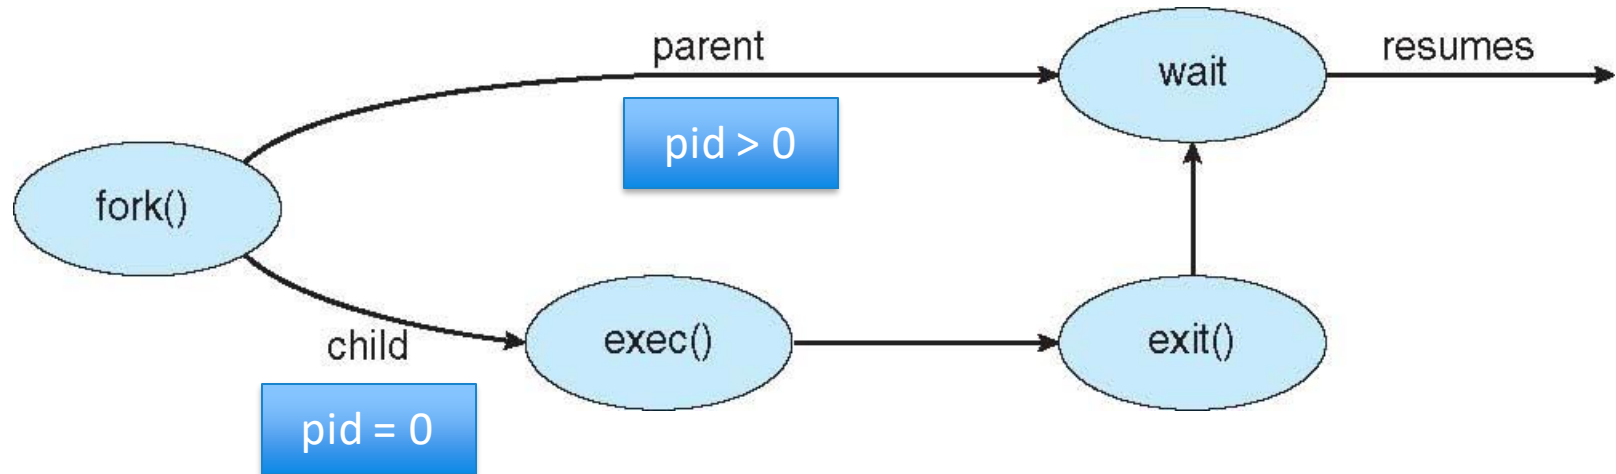
Zombie and orphan process

- Zombie process
 - A process that has terminated, but whose parent has not (yet) called `wait()`
 - All processes move to this state when they terminate, and remain there until parent calls `wait()`
 - Entry in process table removed only after calling `wait()`

Zombie and orphan process

- Zombie process
 - A process that has terminated, but whose parent has not (yet) called `wait()`
 - All processes move to this state when they terminate and remain there until parent calls `wait()`
 - Entry in process table removed only after calling `wait()`
- Orphan process
 - parent terminated without invoking `wait()`
 - Immediately “init” process assigned as parent
 - “init” periodically invokes `wait()`

Usual workflow between parent & child



- Workflow can be different from what is shown
 - E.g., parent may or may not wait for the child
 - E.g., child may or may not use `exec()`

Inter Process Communication

Inter-process communication (IPC)

- Processes executing concurrently in OS may be independent or cooperating
- Cooperating process
 - Affect or be affected by other processes
 - Can share data
 - Speed-up in computation
 - Design can be modular

Inter-process communication (IPC)


- Processes executing concurrently in OS may be independent or cooperating
- Cooperating process
 - Affect or be affected by other processes
 - Can share data
 - Speed-up in computation
 - Design can be modular
- Cooperating processes need IPC
 - Several types of IPC support provided by OS

Inter-process communication (IPC)

- Ways to do IPC
 - 1: shared memory - `shmget()`, `shmcat()`, `shmaddr()`, `shmat()`, `shmdt()`, `shmctl()`
 - 2: message passing (pipe) - `pipe()`, `read()`, `write()`, `close()`
 - 3: message passing (named pipe) - `mkfifo()`, `read()`, `write()`, `close()`
 - 4: over network - RPC or Remote Procedure Call, sockets

Shared memory system

Schematic for shared memory



Require the communicating processes to establish a region of shared memory

Relevant system calls in Linux / Unix
shmget()
shmat()
shmdt()
shmctl()

Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

```
.
```

```
.
```

```
shmdt(myseg);
```

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0);
```

```
.
```

```
.
```

```
shmdt(myseg);
```

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

Create shared memory segment

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

```
myseg = shmat(shmid, NULL, 0)
```

Attach the memory segment to
the address space of this process

```
.
```

```
.
```

```
shmdt(myseg);
```

```
.
```

```
.
```

```
shmctl(shmid, IPC_RMID, NULL);
```

Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

Create shared memory segment

```
myseg = shmat(shmid, NULL, 0);
```

Attach the memory segment to the address space of this process

.

.

```
shmdt(myseg);
```

Detach the memory segment from the address space of this process

.

.

```
shmctl(shmid, IPC_RMID, NULL);
```

Let's check the function calls

```
char *myseg;
```

```
key_t key; int shmid;
```

```
key = 235; // some unique id
```

```
shmid = shmget(key, 250, IPC_CREAT | 0666);
```

Create shared memory segment

```
myseg = shmat(shmid, NULL, 0);
```

Attach the memory segment to the address space of this process

.

.

```
shmdt(myseg);
```

Detach the memory segment from the address space of this process

.

.

```
shmctl(shmid, IPC_RMID, NULL);
```

Mark the segment to be destroyed

Example: Producer consumer problem

- A producer process produces information that is consumed by the consumer process
 - Compiler produces assembly code consumed by assembler
 - Program produces lines to print, print spool consumes
 - The information is written to / read from a buffer

Example: Producer consumer problem

- A producer process produces information that is consumed by the consumer process
 - Compiler produces assembly code consumed by assembler
 - Program produces lines to print, print spool consumes
 - The information is written to / read from a buffer
- Two variants
 - Bounded buffer: producer waits when buffer is full, consumer waits when buffer is empty
 - Unbounded buffer: consumer waits when buffer is empty

Producer consumer solution with bounded buffer

- Shared data: implemented as a circular array

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    . . . // info to be shared
```


```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

These variables reside
in shared memory




Key ideas

- Circular buffer
 - Index in: the next free position to write to
 - Index out: the next filled position to read from
- To check buffer full or empty:
 - Buffer empty: $in == out$
 - Buffer full: $in + 1 \% BUFFER_SIZE == out$
 - This scheme allows at most $BUFFER_SIZE - 1$ items in the buffer

Pseudo code

```
while (true) {  
    /* Produce an item */  
  
    while (( (in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing -- no free buffers */  
  
    buffer[in] = newProducedItem;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

Producer



```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    itemToConsume = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    return itemToConsume;  
}
```

Consumer

Solution is correct, but can only use
BUFFER_SIZE-1 elements

Better utilization of buffer space

- Circular buffer (shared memory)
- Suppose that we want to use all buffer space:
 - an integer count: the number of filled buffers (shared memory)
 - Initially, count is set to 0.
 - incremented by producer after it produces a new buffer
 - decremented by consumer after it consumes a buffer.

Better utilization of buffer space: Pseudo code

Producer

```
while (true) {  
    /* produce an item  
       and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Consumer

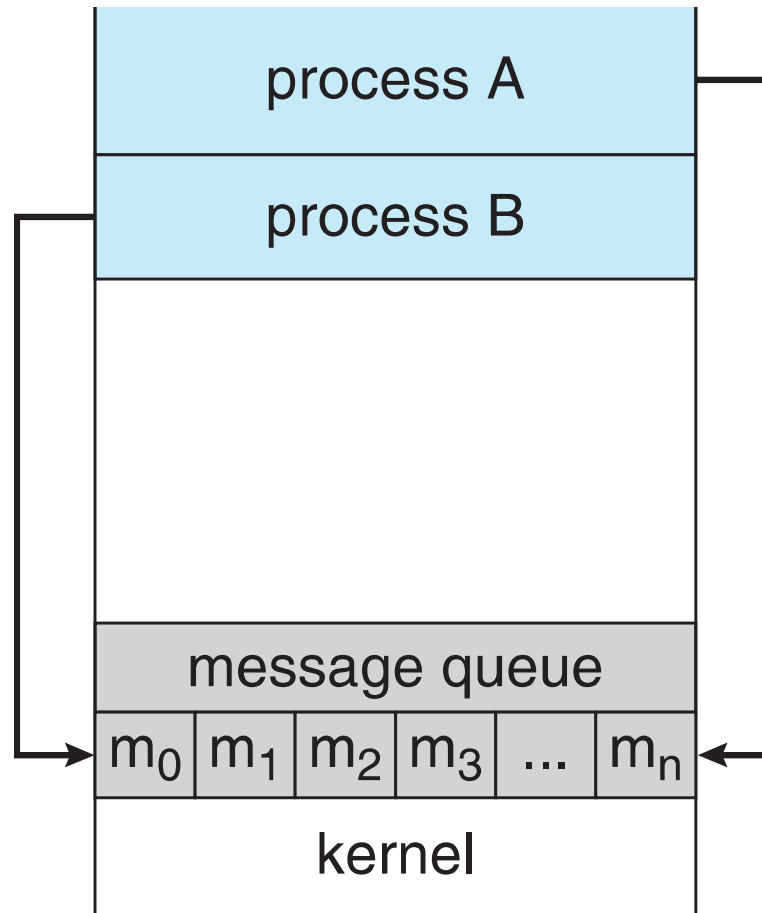
```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in  
       nextConsumed */  
}
```

Message passing system

Basics of message passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*)
 - **receive**(*message*)
- The *message* size is either fixed or variable

Communication model



Ways for message passing

- Pipes
- Named pipes

Pipes

- Acts as a medium to allow two processes to communicate
 - Communication can be uni/bi-directional
 - Must there exist a relationship (i.e., parent-child) between the communicating processes?
 - Can the pipes be used over a network?

Ordinary pipes

- A message passing medium between related processes
 - Cannot be accessed from outside the process
 - Typically, a parent process creates a pipe and uses it to communicate with its child process
 - Pipes behave like FIFO queues
 - Read-write in pipe == producer-consumer
 - Producer writes to one end (the write-end of pipe)
 - Consumer reads from the other end (the read-end of pipe)
 - Unidirectional

Producer consumer in pipes

Producer

```
message next_produced;  
  
while(TRUE){  
    <produce and put data in next produced>  
    ...  
    send( next_produced);  
}
```

Consumer

```
message next_consumed;  
  
while(TRUE){  
    receive(next_consumed);  
    ...  
    <data in next consumed>  
}
```

Named pipes

- Accessed as files by processes
 - No parent-child relation is necessary
 - Still behave like FIFO queues (even called fifo)
 - Several processes can use the named pipe

Named pipes

- Accessed as files by processes
 - No parent-child relation is necessary
 - Still behave like FIFO queues (even called fifo)
 - Several processes can use the named pipe

```
char* myfifo = '/tmp/myfifo';  
mkfifo (myfifo, 0666); // creates the fifo or named pipe (a special file)  
...  
fd = open(myfifo, O_WRONLY); // Process A  
write(fd, ...); // Process A  
close(fd); // Process A  
...  
fd = open(myfifo, O_RDONLY); // Process B  
read(fd, ...); // Process B  
close(fd); // Process B
```


Named pipes

- Once you have created a FIFO special file / named pipe, any process can open it for reading or writing, in the same way as an ordinary file.
- However, it has to be open at both ends simultaneously before you can proceed to do any input or output operations on it.

Finally, for communication of two processes over network

- Sockets API
- Remote procedure call

Summary

- What is a process?
 - Structure of a process
 - Process states
 - Process control block
 - Context switch
- Queues for process scheduling
 - Ready queues, event queues, queueing diagram
- How does two processes talk?
 - Shared memory, pipe, named pipe